

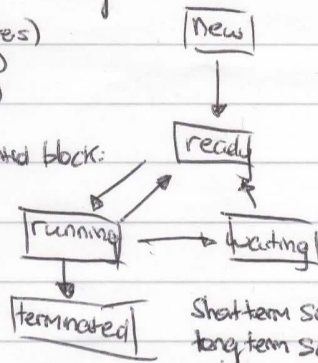
System calls low-level function for common operations  
 Kernel Space where kernel executes and gives services  
 User Space where user processes run

Program passive entity stored on disk as executable  
 Process active, loaded into memory:

- Program code
- Current activity (PC/registers)
- Stack (temp variables)
- Data (global variables)
- heap

OS tracks process with process control block:

- program counter
- process number
- process state
- registers
- memory info
- I/O status



Short-term Scheduler  
 long-term Scheduler  
 dispatcher

Choose what to execute next  
 Chooses what to put in ready queue  
 gives control of CPU to process chosen by ST Scheduler

Job Queue Set of all processes

Ready queue processes ready for execution

Device queue waiting for I/O devices

Independent process cannot affect other process  
 but cooperating process can  
 good for information sharing / computational  
 Speedup | modularity | convenience

Data parallelism Subsets of same data operated on in different cores

Task parallelism different tasks performed on different cores

Amdahl's law Speedup with more cores  $\approx \frac{S}{1 + \frac{1-S}{N}}$

Implicit threading When compilers and runtime libraries control threading

Four instances that trigger scheduler

running  $\rightarrow$  waiting  
 running  $\rightarrow$  ready  
 waiting  $\rightarrow$  ready  
 terminates

Race condition when two  
 processes write to same variable  
 but only one value is  
 preserved

Starvation very low priority processes  
 never get executed

Aging gradually increase priority  
 of waiting processes

Asymmetric multiprocessing  
 one processor does scheduling while  
 others run code

Symmetric multiprocessing  
 each manages own scheduling

Safe state if no deadlocks

Unsafe state if possible deadlocks

Priority inversion, the lower priority  
 process is secretly a higher  
 priority than that assigned  
 to it

Simple Structure System Progs  $\rightarrow$  OS Drivers  
 Application Progs.  $\rightarrow$  BIOS Drivers  
 Difficult to maintain  
 Little overhead sys calls

Layered Approach

Bottom layer hardware, top layer user interface, clear interface  
 between layers, easy to construct and debug, hard to  
 define layers, inefficient with lots of system calls

Microkernel

Non-essential programs in user mode, essential in kernel mode,  
 extending OS is easy, increased system call overhead

Modular

Extends upon microkernels, kernel modules loaded when  
 needed, do not need to rebuild to add functionality

I/O bound does I/O more, CPU bound does computation  
 more

Context Switch when CPU changes to another process

1. Save State of current process in process control block
2. load saved State of new process

Shared Memory

processes share a region of memory  
 only need system calls for setup  
 have to avoid conflicts/race conditions

Message Passing

processes send message to each other  
 send(m) and receive(m)  
 must establish a connection  
 direct communication (name each other)  
 indirect communication (mailboxes)  
 blocking (sync), nonblocking (async)

Thread is

basic unit of CPU execution  
 contains thread id, program counter, register set, stack  
 threads are cheaper to create, good for:

responsiveness  
 resource sharing  
 economy  
 scalability

threadpool is where  
 they await for work

nonpreemptive no choice  
 preemptive choices

CPU utilisation  
 throughput

keep the CPU as busy as possible  
 number of processes that complete their execution per time unit

turnaround time

amount of time to execute a particular process

waiting time

amount of time a process has been in waiting queue

response time

time from request to first response

Critical Section Problem

1. only one can be in it's critical section at any time
2. if none in critical section and some waiting, one go in
3. can only wait a limited number of times before going in

Mutex lock

Process must acquire lock before entering  
 critical section and release when done  
 busy-waiting

Deadlock occurs when two or more processes wait indefinitely for an event  
 that can only be caused by the waiting processes

Deadlock occurs if four conditions hold:

1. if only one process can use resource at a time
2. a process holding a resource is waiting for a resource held by another
3. a resource can only be released after a process finishes its tasks
4.  $P_0, P_1, \dots, P_{n-1}$ ,  $P_0$  waiting for  $P_1$  res,  $P_1$  for  $P_2$  res, ...,  $P_{n-1}$  for  $P_0$  res

external fragmentation if total memory space exists to satisfy request  
 but not contiguous / internal fragmentation if allocated memory  
 is slightly larger than requested memory

Semaphore

Semaphore is an integer  
 process puts itself in wait state until can run  
 control access to a finite number of a  
 given resource

Preemptive Kernel allows a  
 process to be preempted  
 while running in kernel mode  
 Nonpreemptive kernel does  
 not allow preempting